

# AI AND IOS

Aileen Nielsen

[aileen.a.nielsen@gmail.com](mailto:aileen.a.nielsen@gmail.com)

## WHAT I WAS GOING TO TALK ABOUT

Pythonic aspects of Swift

Machine learning with Swift

Building neural networks for your apps

**But then...WWDC 2017 & CoreML**

## OUTLINE

Core ML & related  
Machine learning concepts  
Final bits and bobs

**Goal: empowerment**

# CORE ML & RELATED

## BEFORE WWDC 2016

- **Data scientists don't always know much about the implementation details** of the techniques they use
  - Especially bad/poorly trained data scientists
  - They may know about 'neural networks' but don't know how to break their models up into component pieces for implementation
- **But** iOS engineers are iOS engineers, **not data engineers or data scientists**
- **But** many interesting new applications are going to use some form of intelligence...and this trend will only accelerate
- **Open source projects** to
  - Translate existing ML models into code
  - Train models
  - Facilitate data-sciencey stuff



**Swift A.I.**  
A.I. LIBRARY FOR SWIFT

### MACHINE Learning

Machine Learning for the Mac

#### Intro

MACHINE Learning (pron. 'maek[ən] 'la:rnɪŋ) is framework that provides a quick and easy way to experiment Machine Learning with native code on the Mac, with some specific support for Natural Language Processing. It is written in Objective-C, but it is compatible by Swift.

Currently the framework supports:

- [Neural Networks][#Neural Networks]
- [Bag of Words][#Bag of Words]
- [Word Vectors][#Word Vectors]

(swix)

### Swift Matrix and Machine Learning Library

Apple's Swift is a high level language that's asking for some numerical library to perform computation *fast* or at the very least *easily*. This is a bare-bones wrapper for that library.

A way to have iOS run high-level code similar to Python or Matlab is something I've been waiting for, and am incredibly excited to see the results. This will make porting complex signal processing algorithms to C *much* easier. Porting from Python/MATLAB to C was (and is) a pain in the butt, and this library aims to make the conversion between a Python/Matlab algorithm and a mobile app *simple*.

WWDC 2016

## WWDC 2016

**BNNS** introduced a set of routines in the **Accelerate (CPU) framework** designed to implement neural networks, introduced in iOS 10

- Supports only 3 types of layers: Convolution, Fully Connected, Pooling

**MPSCNN** introduced a set of routines with **Metal (GPU) framework** to implement neural networks, introduced in iOS 10

- Also build neural networks **layer by layer**
- **More opportunities for customization** than BNNS
- Similar but not identical API

**WWDC Emphasis on AI:** speech recognition, user intent, SiriKit, machine learning, and NNs

# WWDC 2017

## CoreML

- Close to a **'set it and forget it'** style API
- Uses Metal & Accelerate under the hood

## Neural-network-related enhancements to **Metal**

- **More kinds of kernels**
- API for **building graphs** makes this **more Keras like**
- Kernels are now **serializable**
- **More datatypes**, in particular the preferred 16-bit float is now supported

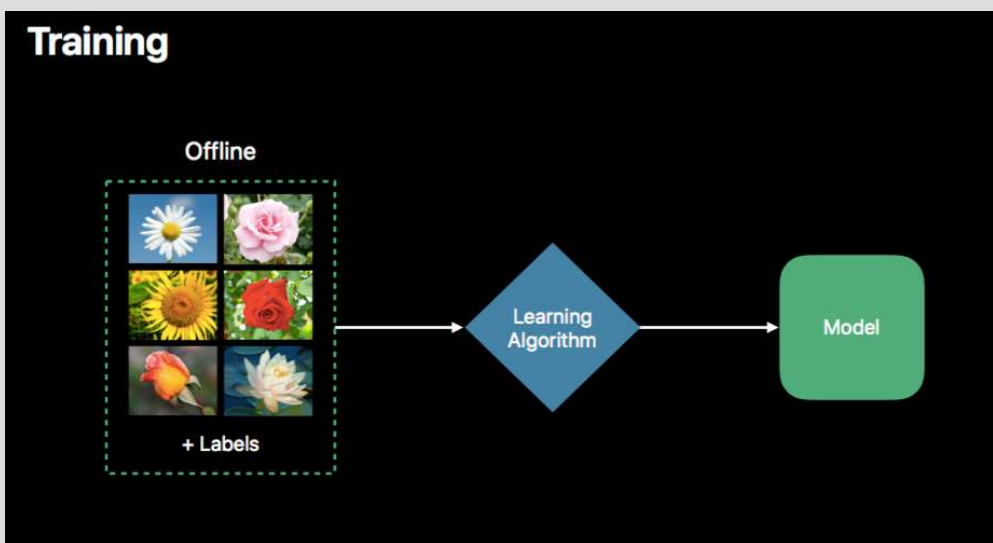
Seems **BNNS** did not get much updating



# CORE ML IN 5 MINUTES



# CORE ML IN 5 MINUTES



# CORE ML IN 5 MINUTES

## Inference



Model



Label: **Rose**  
Confidence: **95%**

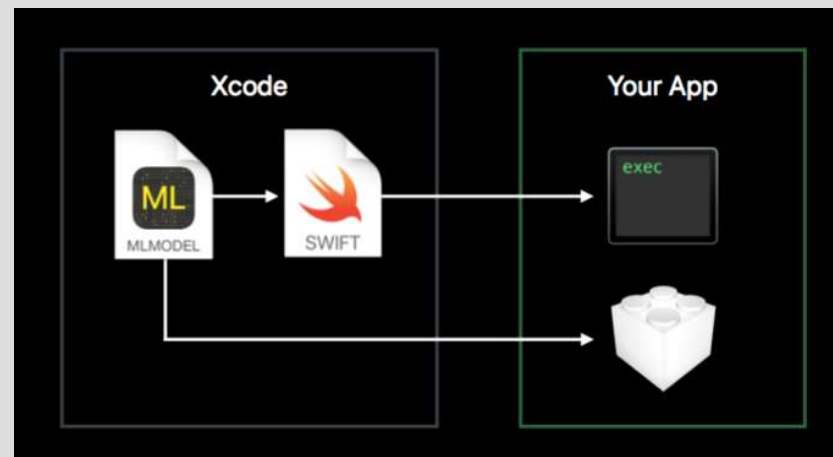
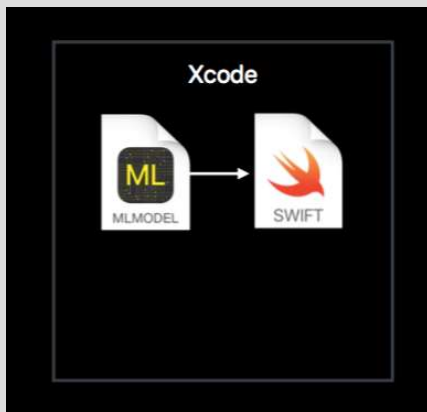
# CORE ML IN 5 MINUTES



```
import coremltools

keras_model = 'flowers.h5'
model = coremltools.converters.keras.convert(
    keras_model,
    image_input_names = 'data',
    class_labels = 'labels.txt')

model.save('FlowerClassifier.mlmodel')
```



# CORE ML IN 5 MINUTES

## Machine Learning Model

Name: FlowerClassifier  
Type: Neural Network Classifier  
Size: 41.6 MB  
Author: Lizi Ottens  
License: MIT  
Description: Identify the type of flower present in an image.

## Model Class

FlowerClassifier (Swift generated source)

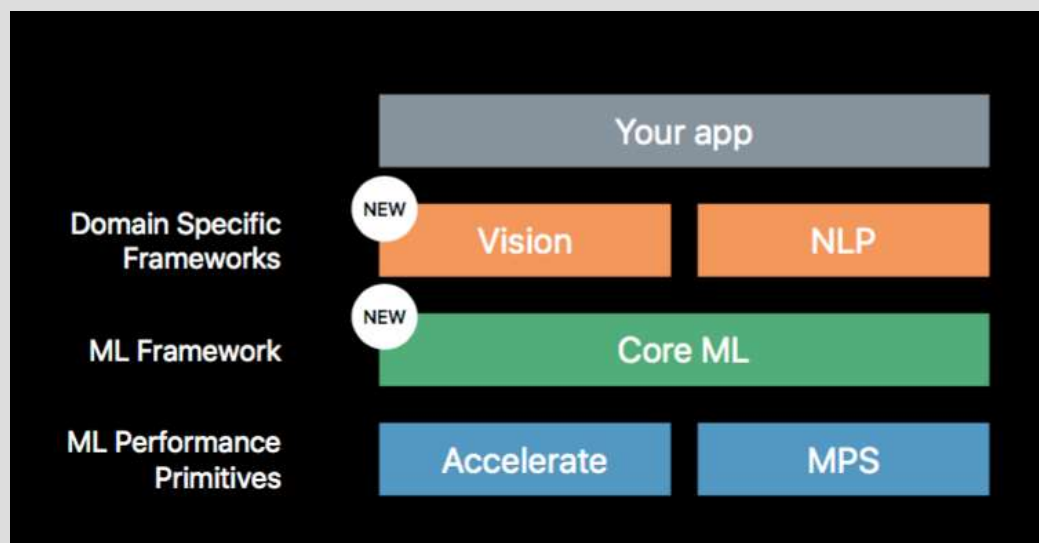
## Model Evaluation Parameters

Name	Type	Description
Inputs		
flowerImage	Image<RGB,227,227>	Input image of a flower
Outputs		
flowerType	String	Most likely flower type in image
flowerTypeProbs	Dictionary<String,Double>	Probability of each flower type

```
class FlowerClassifierInput {  
    var flowerImage: CVPixelBuffer  
}  
  
class FlowerClassifierOutput {  
    let flowerType: String  
    let flowerTypeProbs: [String: Double]  
}  
  
class FlowerClassifier {  
    convenience init()  
    func prediction(flowerImage: CVPixelBuffer) throws -> FlowerClassifierOutput  
}
```

```
let flowerModel = FlowerClassifier()  
if let prediction = try? flowerModel.prediction(flowerImage: image) {  
    return prediction.flowerType  
}
```

# CORE ML IN 5 MINUTES



## COOL FEATURES

CoreML **decides which parts of a model to run on CPU vs. GPU** (so even the conscientious iOS engineer does not need to run CPU vs. GPU variants to see which is more performant)

- GPU is best for many parallel identical operations
- CPU is good for memory intensive or computationally serialized processes
- In general GPU tends to be better

Can run it from the **simulator**

- (unlike Metal)

In subsequent beta versions, got possibility to download and compile **mlmodel** so you don't have to release a new app to update your model

- This is **amazing** if your app basically is your model

## CURRENTLY SUPPORTED MODELS

- support vector machines (SVM)
- tree ensembles such as random forests and boosted trees
- linear regression and logistic regression
- neural networks: feed-forward, convolutional, recurrent

**More on these later...**



## DOWNSIDERS TO COREML (OR OPPORTUNITIES TO LEARN)

- You probably need to **learn some Python**
  - But it's very Swiftly
  - Or rather, Swift is very Pythonic
- **You will be responsible for understanding what is going wrong in models on device**
  - So you'll need to understand the basics of models
- Ultimately this will emerge as a **specialization** for iOS engineers
  - Model translation
  - Model implementation
  - Model performance enhancements and pitfalls
- **Not (yet?) open source**, so it will always be behind the data science tools

## CORE ML GOTCHAS

Models can **run fine on simulator but crash on device**

- E.g. on **simulator model ran on the CPU** and on device ran on **GPU** and did something impermissible

Many early '**bugs**' center around **not giving the right kind of input**

- There are **many ways to describe an image**, but only one that the model expects
- Unexpected model behavior is **usually the result of not matching expected input**

## PREDICTIONS FOR CORE ML

- **Fine tuning**
- **R** support
  - Finance – various packages but **glmnet** and **elasticnet** for starters
  - Sociology/politics – various packages, but **glmnet** and **survey** are possibly most important
  - Biostats – **Bioconductor** package
- **Unsupervised learning**
  - But this is **computationally taxing** so just as likely as on-device learning
  - **Everything** is a lot less demanding than training neural networks

# ML CONCEPTS




# REGULARIZATION

$$Obj(\Theta) = L(\Theta) + \Omega(\Theta)$$

**Training Loss** measures how well model fit on training data

**Regularization**, measures complexity of model


# OVERFITTING (AND UNDERFITTING)



# OVERFITTING DIAGNOSIS

Error

Cross-validation error



## OTHER TERMS YOU'LL ENCOUNTER

- Preprocessing
- Cross-validation
- Test & training sets
- Wide data vs. long data
- Normalized data
- Model compression





## CURRENTLY SUPPORTED MODELS

- **linear regression** and logistic regression
- support vector machines (**SVM**)
- **tree ensembles** such as random forests and boosted trees
- **neural networks**: feed-forward, convolutional, recurrent

# REGRESSION MODELS

- Linear regression models try to **predict a dependent value via a linear relationship** with a set of predictors



## REGRESSION MODELS

- Linear regression models try to **predict a dependent value via a linear relationship** with a set of predictors
  - Normal methodology makes vanilla **assumptions about input data**
  - However **assumptions aren't always true** in the real world but there is **usually a workaround**
  - **Computationally cheap** relative to many other methods
  - Remains workhorse of prediction
- **Logistic regression** uses linear methods but **seeks to model a probability**, which is bounded, and hence requires a linear separation
- Most real-world models will use some form of **regularization**
  - 'wide' datasets

## PROS AND CONS OF REGRESSION MODELS



### **Pros**

- Straightforward statistical models for generating confidence intervals Logistic regression provides probabilities, so we get classification + probability instead of just classification
- Some closed form solutions
- Well developed refinements for tricky data problems

### **Cons**

- We don't always have a good reason to think linear relationships are linear
- Input data is not always independent even though it is assumed to be so
- Sensitive to outliers
- Only looks at the mean of a dependent variable without looking at the distribution

# SUPPORT VECTOR MACHINES



# SUPPORT VECTOR MACHINES

- The support vector machine is a **separating plane** which best separates groups
- SVM selects a hyper plane which
  - First **maximizes classification accuracy**
  - Second **maximizes minimum distance from hyperplane** to members of each group (**maximizing the margin**)
- Possible to have a **non-linear hyperplane**
  - But usually want to **justify decision** to have a non-linear hyperplane
  - You can fit any set of points with enough parameters, that's a danger called **overfitting**

## PROS AND CONS OF SVM

### Pros

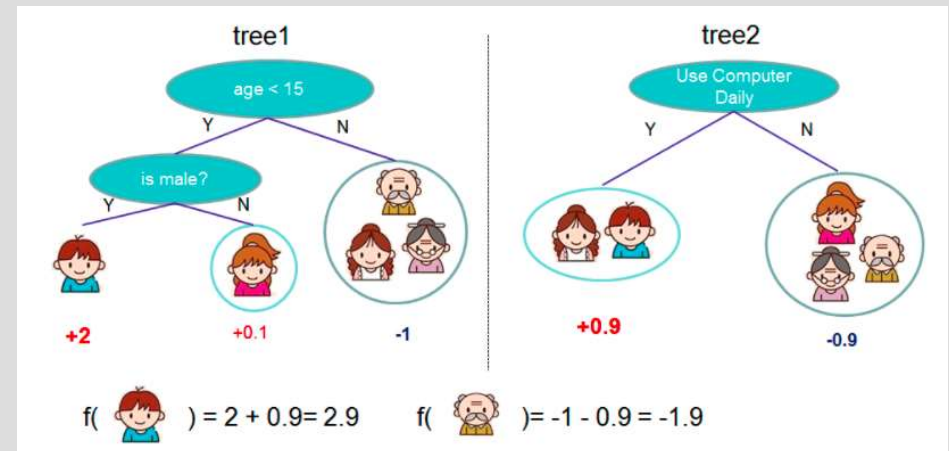
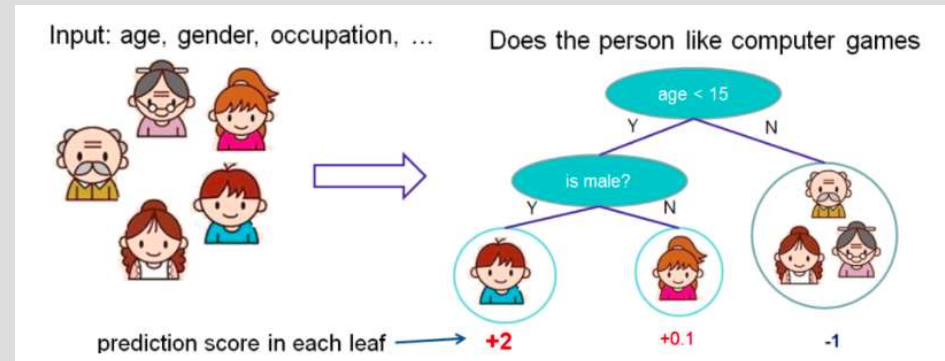
- Relatively **memory efficient** so can be used in high dimensional spaces
- Works when **number of dimensions is greater than number of samples**
- Works really well with **strong signal (large margin)**
- More **robust to outliers** than regression

### Cons

- Does not work so well with **noisy sampling** (small margin)
- With larger data set, very **long training times**
- Does **not provide probability estimates**

# TREE ENSEMBLES

- **Trees** are a fairly intuitive way of making both categorical and numerical predictions
- **Random forests:** It turns out that a bunch of 'dumb' trees together can do a lot better than a single tree





## PROS AND CONS OF TREE ENSEMBLES


### **Pros**

- Random ensemble methods are robust to highly correlated input variables
- High accuracy with uncomplicated methodology


### **Cons**

- Ensembles are not easy to visually interpret

# NEURAL NETWORKS



# CONVOLUTIONAL NEURAL NETWORKS



Original image




Center of Circle Blur on the right side of crab

<https://cambridgespark.com/content/tutorials/convolutional-neural-networks-with-keras/index.html>


<https://www.cs.nyu.edu/~fergus/papers/zeilerECCV2014.pdf>

<https://documentation.apple.com/en/motion/usermanual/index.html#chapter=16%26section=9%26tasks=true>



# CONVOLUTIONAL NEURAL NETWORKS




# RECURRENT NEURAL NETWORK



# RECURRENT NEURAL NETWORK



# STATE OF THE ART ...ISH INCEPTION V. 3



FINAL BITS AND BOBS



## UPSIDES OF CORE ML ACCORDING TO APPLE

Apple listed the benefits of running on device as

- User **privacy**
- **Data** cost
- **Server** cost

But...

- **Larger footprint** on the device
- **Can correct models** on the server-side
- Server-side can **keep the model builders closer to the production** version of the model
- Apart from image-related applications, many models run on **small input data**

## AND SOME DOWNSIDES

- For some organizations, the model is the **core IP** and **core worth** of the enterprise...so organizations may not want to include more engineers on the list of people familiar with the model
- **Model performance/correctness checking may get forgotten**

# ETHICS

- Ethics in AI and machine learning is an **increasingly fraught area**
- Possibility that there is **unintentional racial/gender discrimination**
  - usually done by proxy such as using zip codes which can be a strong proxy for race
- This **may become an increasingly regulated area**
- You have the **obligation to make sure you are putting out correct models**, even if you didn't build them
- Some tension **between privacy and fairness** because how can you check downstream effects?

# PRIVACY

- **Apple is emphasizing privacy** in all its AI implementations
- But at some point **bug reporting** is going to clash with privacy
- The **initial (and corrective) data** has to come from somewhere

## WHAT DOES THIS MEAN FOR YOUR CAREER?

- You should get to talk to **more kinds of people** at your organizations
- As always, the value play is what can you **add** to the APIs Apple is providing?
- New demands on **QA & bug searches**
- **New frontiers**

# THERE ARE STILL INTERESTING THINGS GOING ON BEYOND APPLE API

## DSLR-Quality Photos on Mobile Devices with Deep Convolutional Networks

Andrey Ignatov, Nikolay Kobyshev, Kenneth Vanhoey, Radu Timofte, Luc Van Gool

ETH Zurich

andrey.ignatoff@gmail.com, {nk, vanhoey, timofte, vangool}@vision.ee.ethz.ch

### Abstract

Despite a rapid rise in the quality of built-in smartphone cameras, their physical limitations — small sensor size, compact lenses and the lack of specific hardware, — impede them to achieve the quality results of DSLR cameras. In this work we present an end-to-end deep learning approach that bridges this gap by translating ordinary photos into DSLR-produced images. We propose learning the translation function using a residual convolutional neural network that improves both color rendition and image sharpness.



Figure 1: Sony smartphone image enhanced to DSLR-quality by our method. Best zoomed on screen.

18 Apr 2017

# THERE ARE STILL INTERESTING THINGS GOING ON BEYOND APPLE API

## Real-Time Visual Place Recognition for Personal Localization on a Mobile Device

Michał Nowicki · Jan Wietrzykowski · Piotr Skrzypczyński



**Fig. 4** The trajectories of two PUT MC test sequences: PUT MC short, marked by orange dots, and PUT MC long, marked by blue dashes



**Fig. 5** OpenFABMAP incorrectly recognized locations: A) due to many features placed on windows, B) due to people occluding the view

# THERE ARE STILL INTERESTING THINGS GOING ON BEYOND APPLE API

## Two-view 3D Reconstruction for Food Volume Estimation

Joachim Dehais, *Student Member, IEEE*, Marios Anthimopoulos, *Member, IEEE*, Sergey Shevchik,  
Stavroula Mougiakakou, *Member, IEEE*




Fig. 6. Volume estimation: food surface (separated using the segmentation map), plate surface and vertical projection on the dish plane.



THE HARD THINGS ARE EASY,  
AND **THE EASY THINGS ARE HARD**